

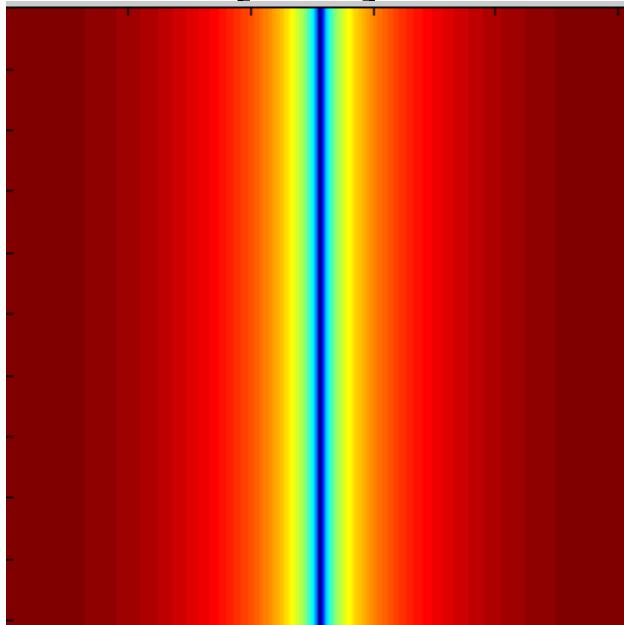
Review of Filtering

- Filtering in frequency domain
 - Can be faster than filtering in spatial domain (for large filters)
 - Can help understand effect of filter
 - Algorithm:
 1. Convert image and filter to fft (fft2 in matlab)
 2. Pointwise-multiply ffts
 3. Convert result to spatial domain with ifft2

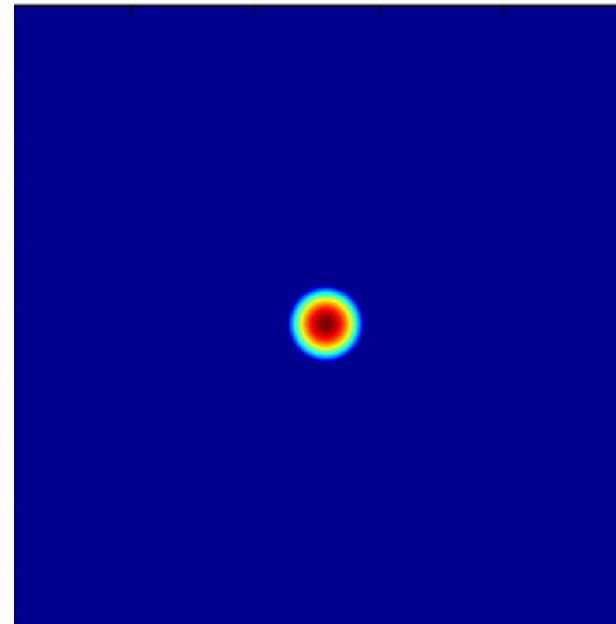
Review of Filtering

- Linear filters for basic processing
 - Edge filter (high-pass)
 - Gaussian filter (low-pass)

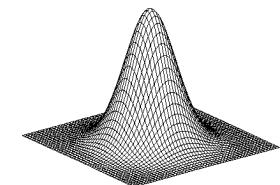
$[-1 \ 1]$



FFT of Gradient Filter

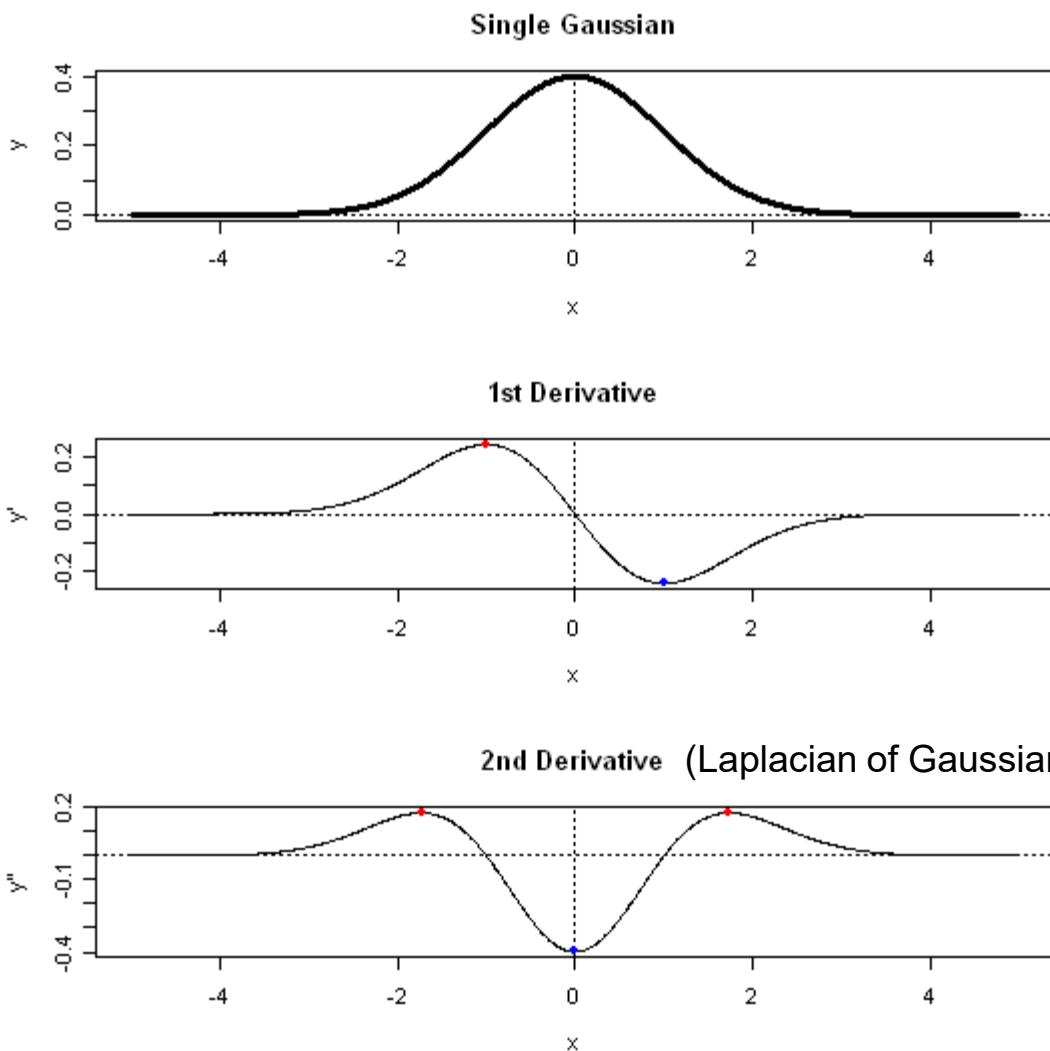


FFT of Gaussian

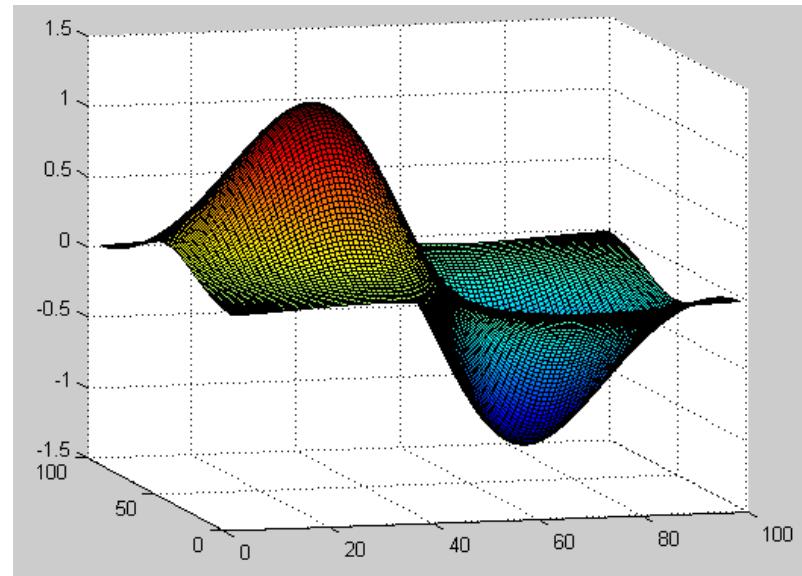


Gaussian

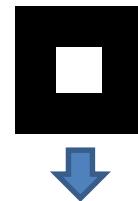
More Useful Filters



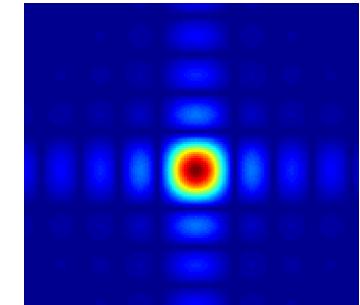
1st Derivative of Gaussian



Things to Remember

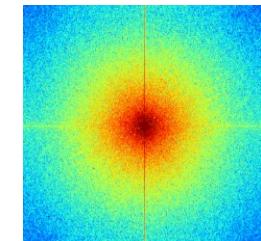


- Sometimes it makes sense to think of images and filtering in the frequency domain
 - Fourier analysis

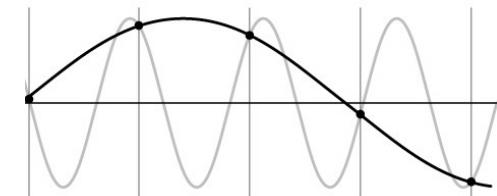


- Can be faster to filter using FFT for large images
 - $N \log N$ vs. N^2 for auto-correlation

- Images are mostly smooth
 - Basis for compression



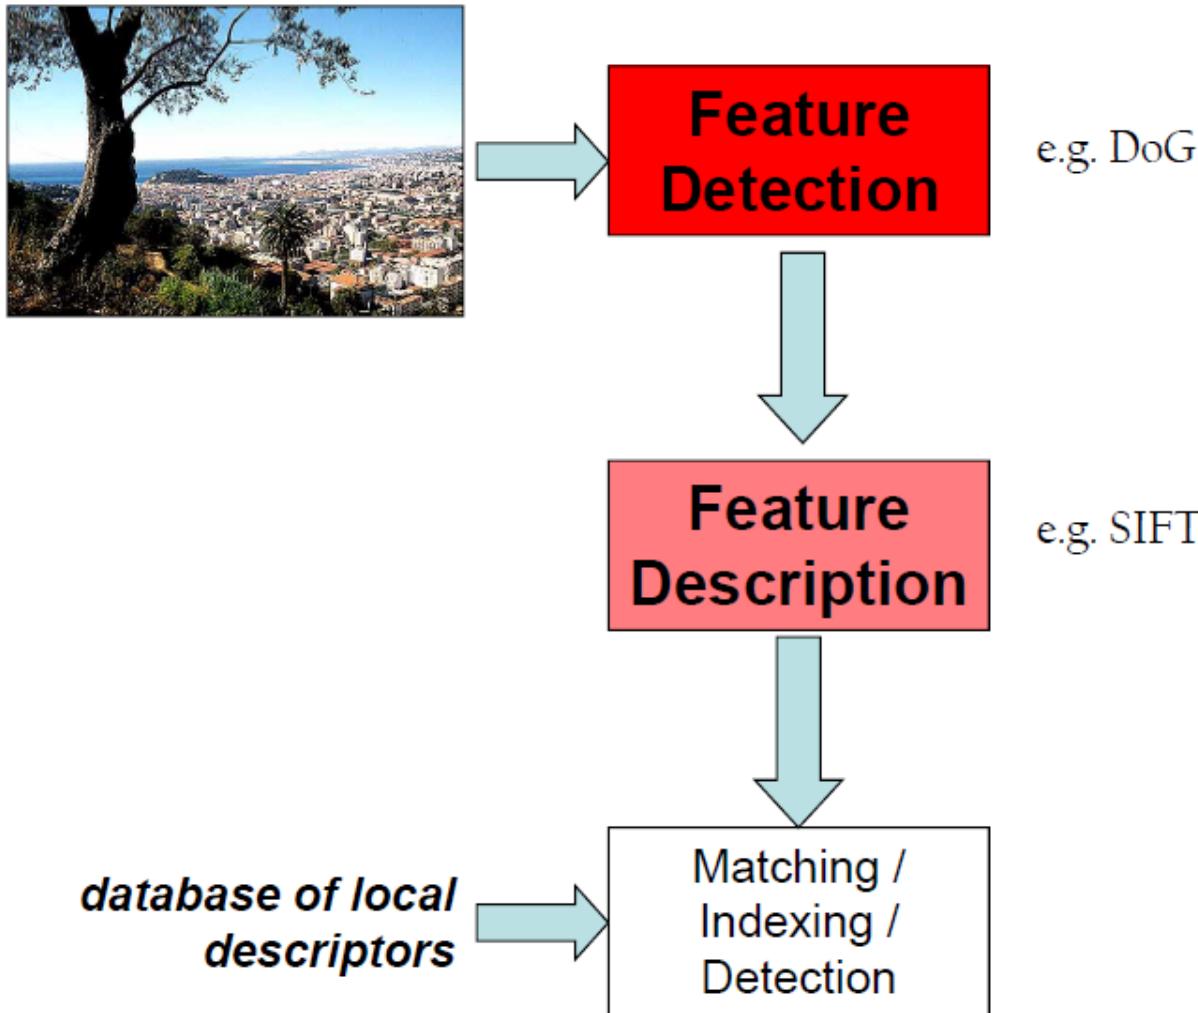
- Remember to low-pass before sampling
 - Otherwise you create aliasing



Previous Lectures

- We've now touched on the first three chapters of Szeliski.
 - 1. Introduction
 - 2. Image Formation
 - 3. Image Processing
- Now we're moving on to
 - 4. Feature Detection and Matching

The big picture...



EDGE / BOUNDARY DETECTION

Szeliski 4.2

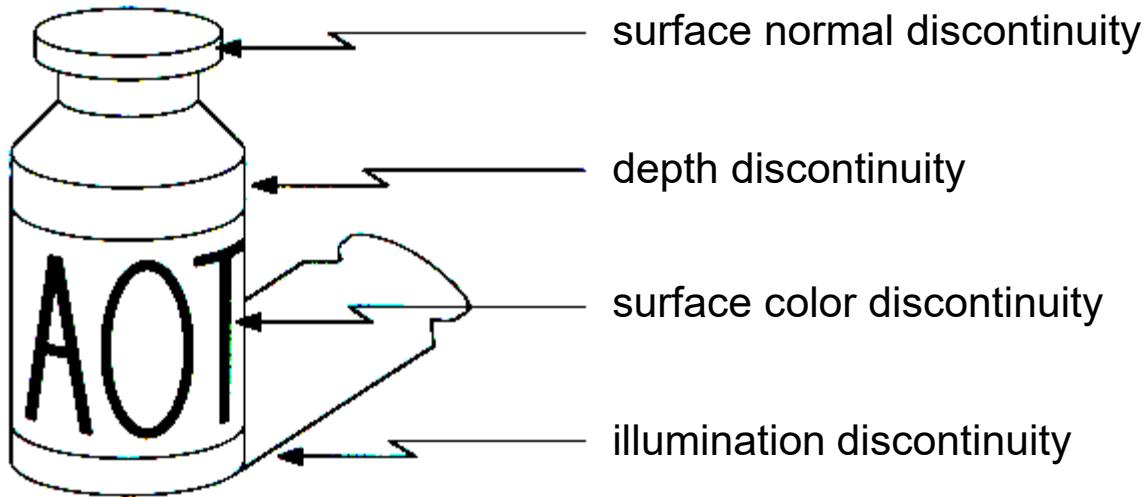
Many slides from James Hays, Lana Lazebnik, Steve Seitz, David Forsyth, David Lowe, Fei-Fei Li, and Derek Hoiem

Edge detection

- **Goal:** Identify visual changes (discontinuities) in an image.
- Intuitively, semantic information is encoded in edges.
- What are some ‘causes’ of visual edges?



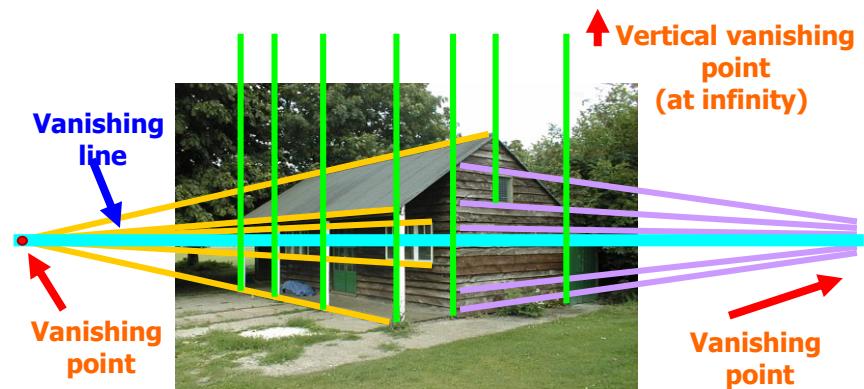
Origin of Edges



- Edges are caused by a variety of factors

Why do we care about edges?

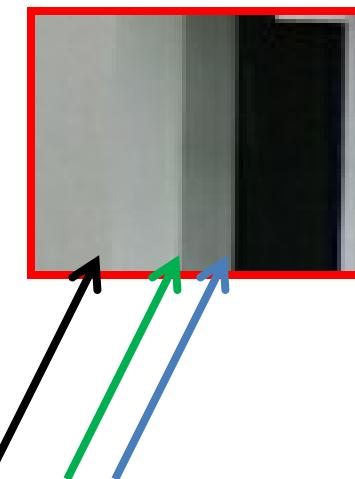
- Extract information
 - Recognize objects
- Help recover geometry and viewpoint



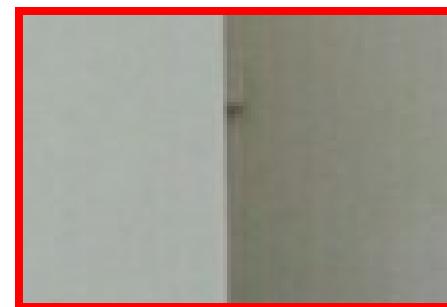
Closeup of edges



Closeup of edges



Closeup of edges

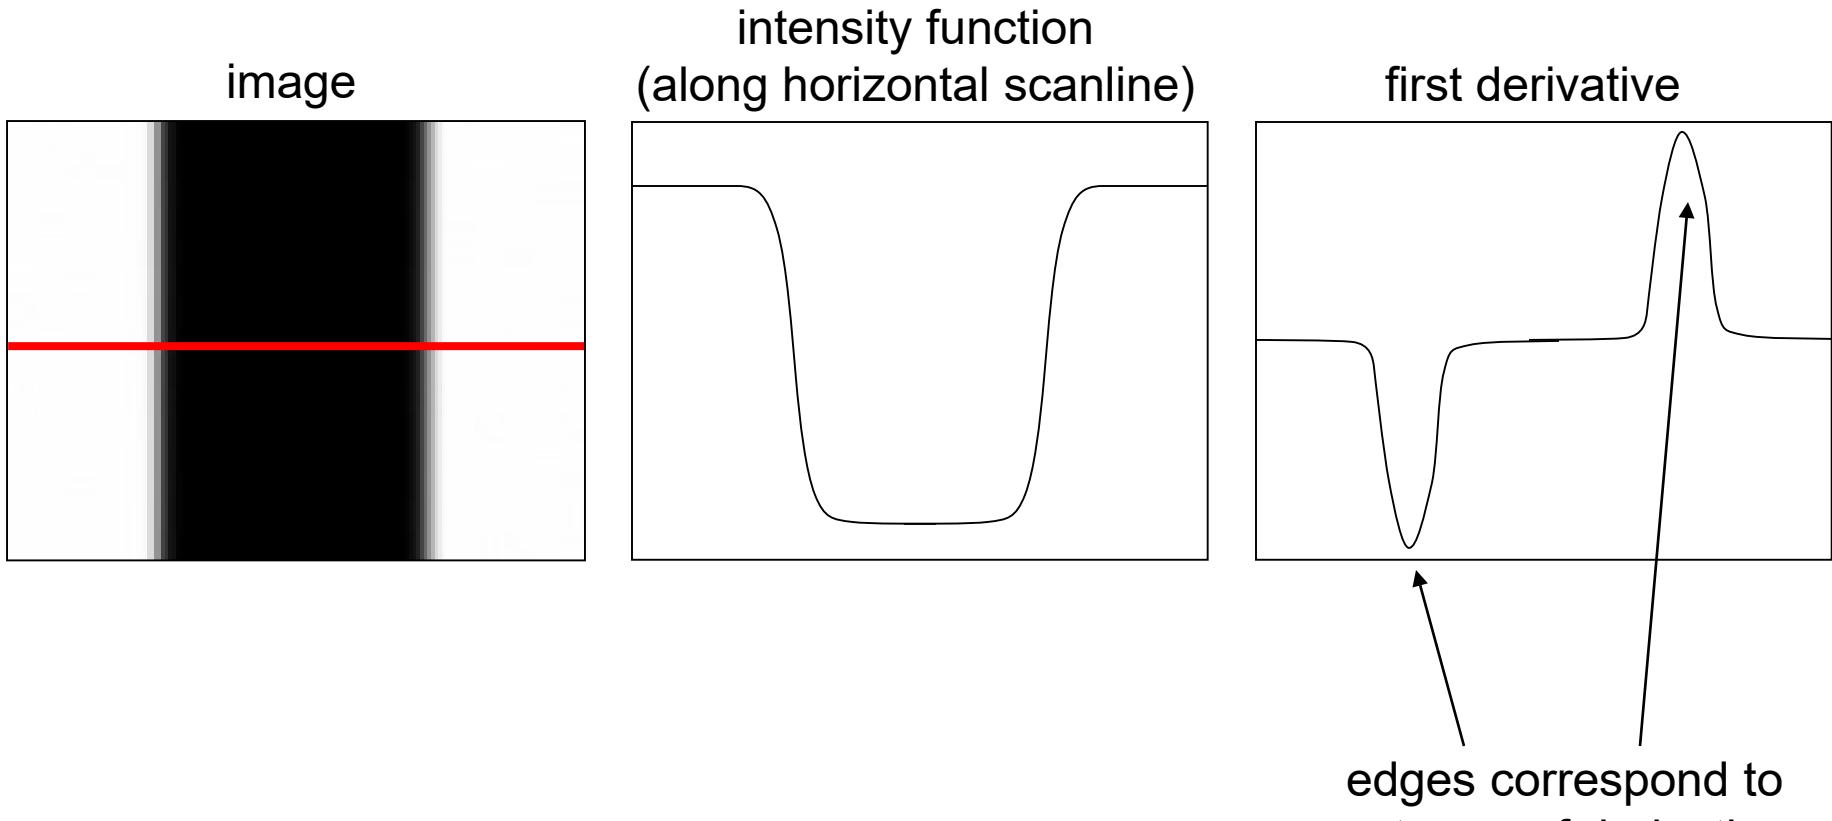


Closeup of edges

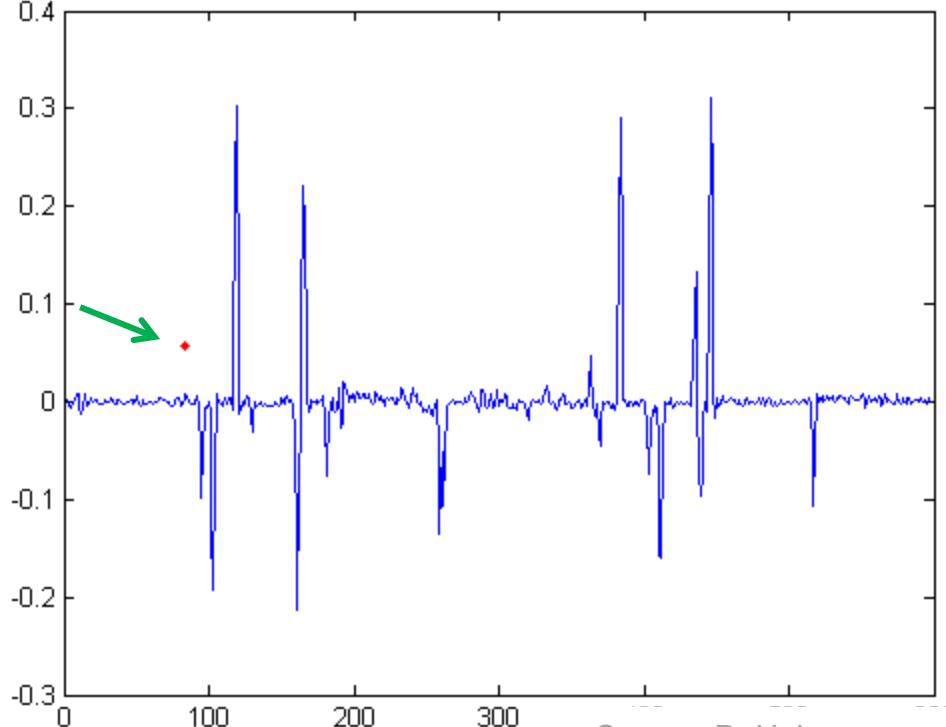
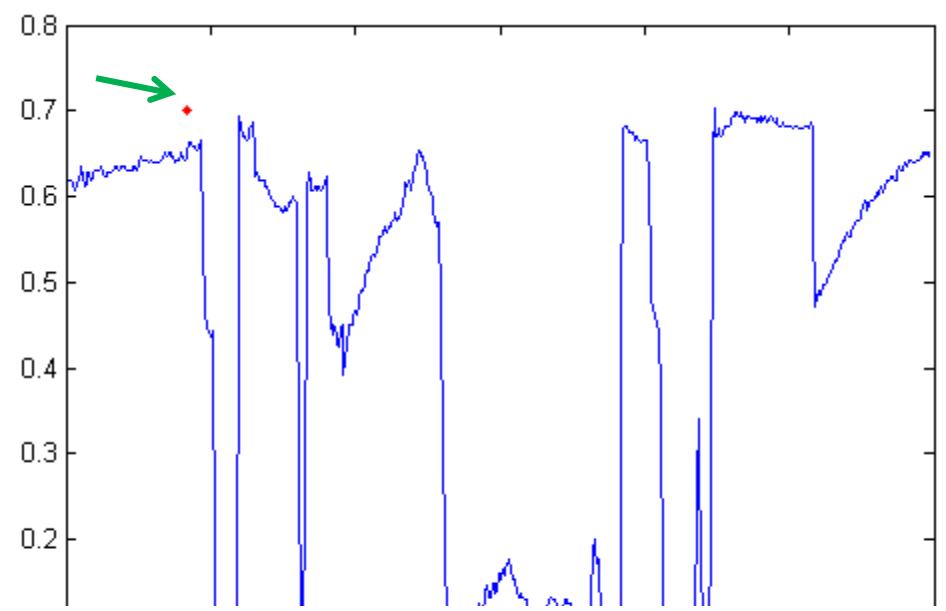
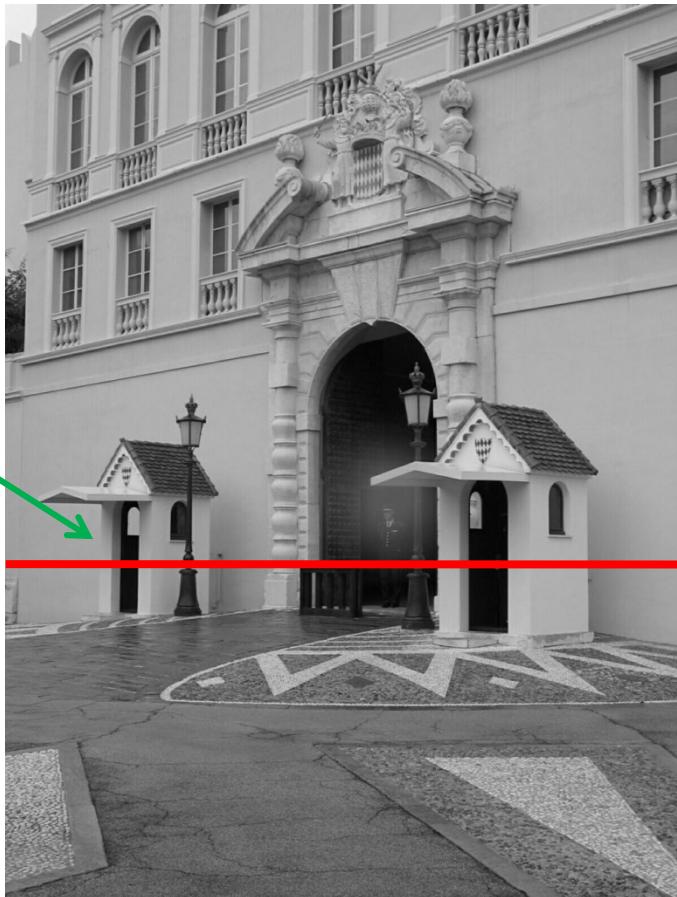


Characterizing edges

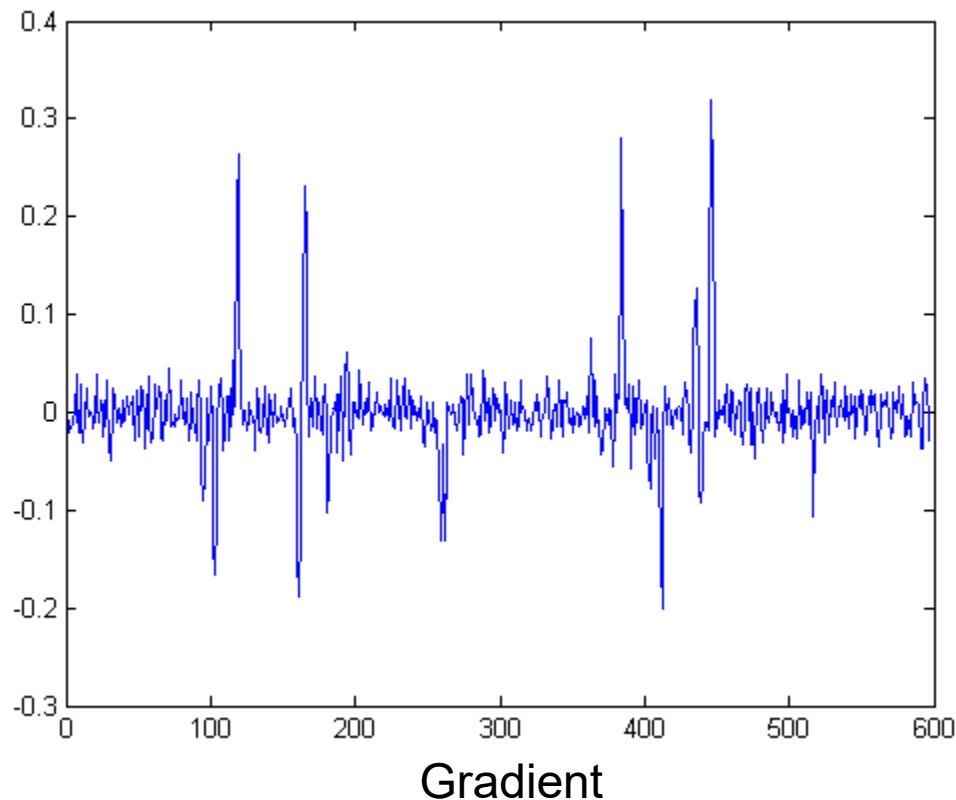
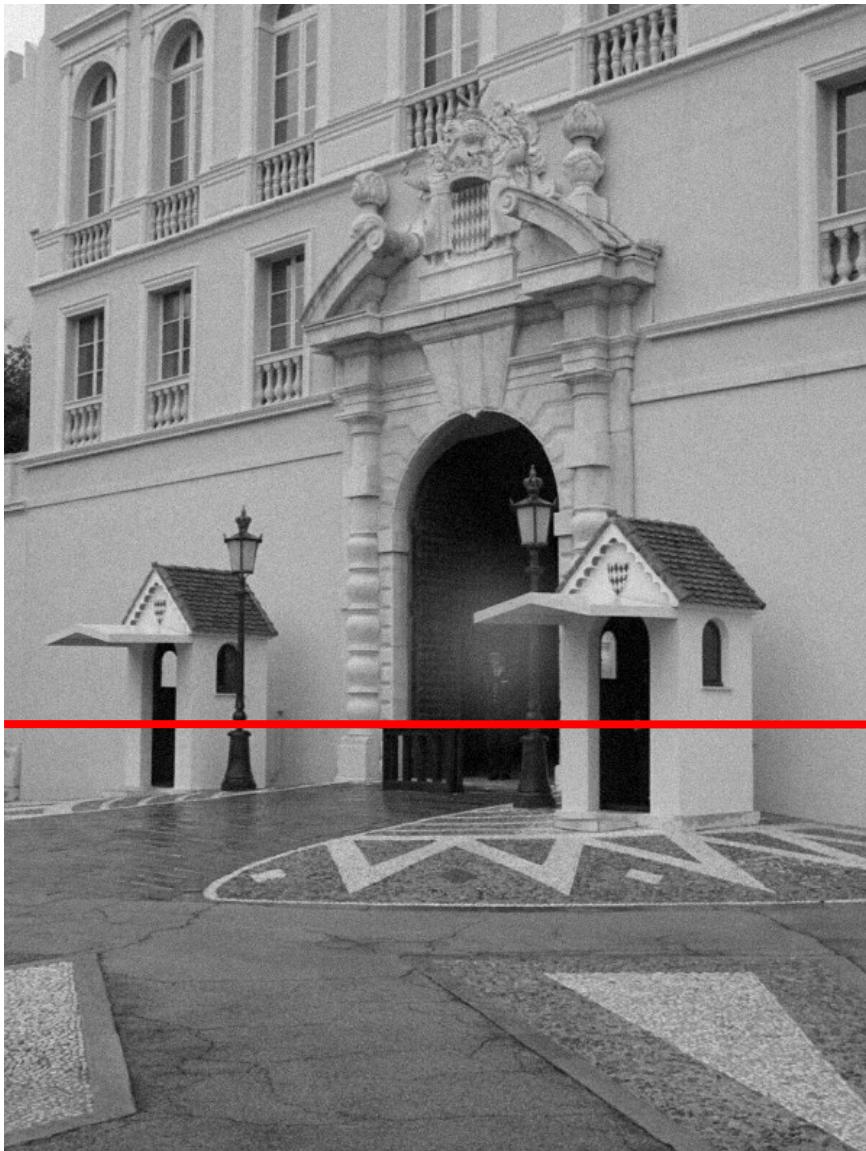
- An edge is **a place of rapid change in the image intensity function**



Intensity profile



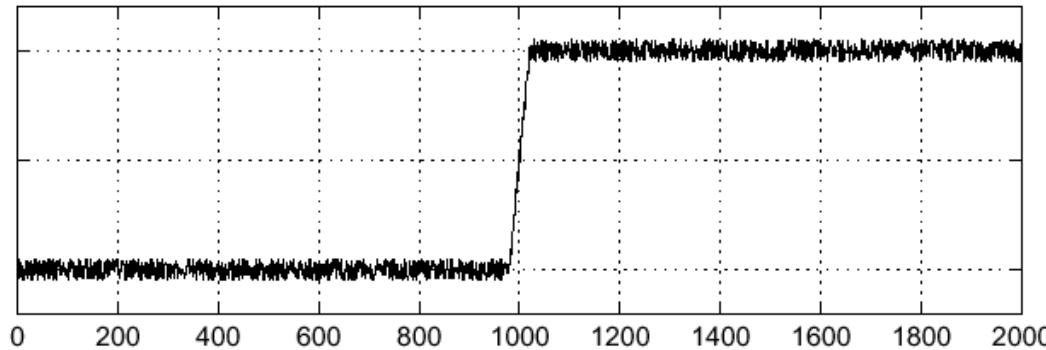
With a little Gaussian noise



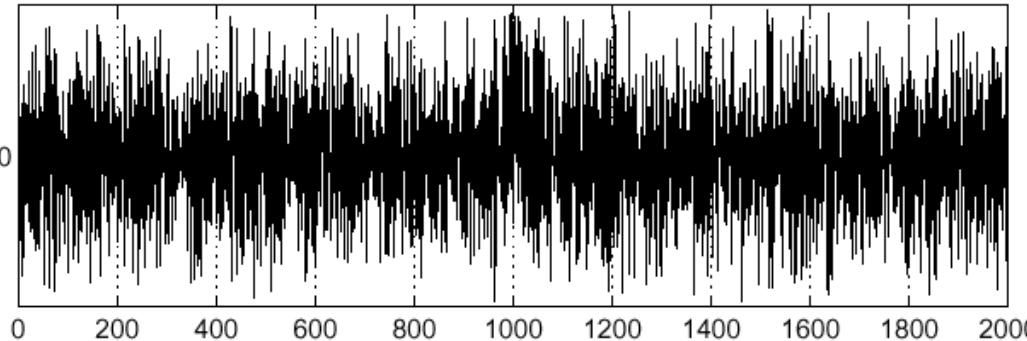
Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal

$f(x)$



$\frac{d}{dx}f(x)_0$



Where is the edge?



未来媒体研究中心
CENTER FOR FUTURE MEDIA



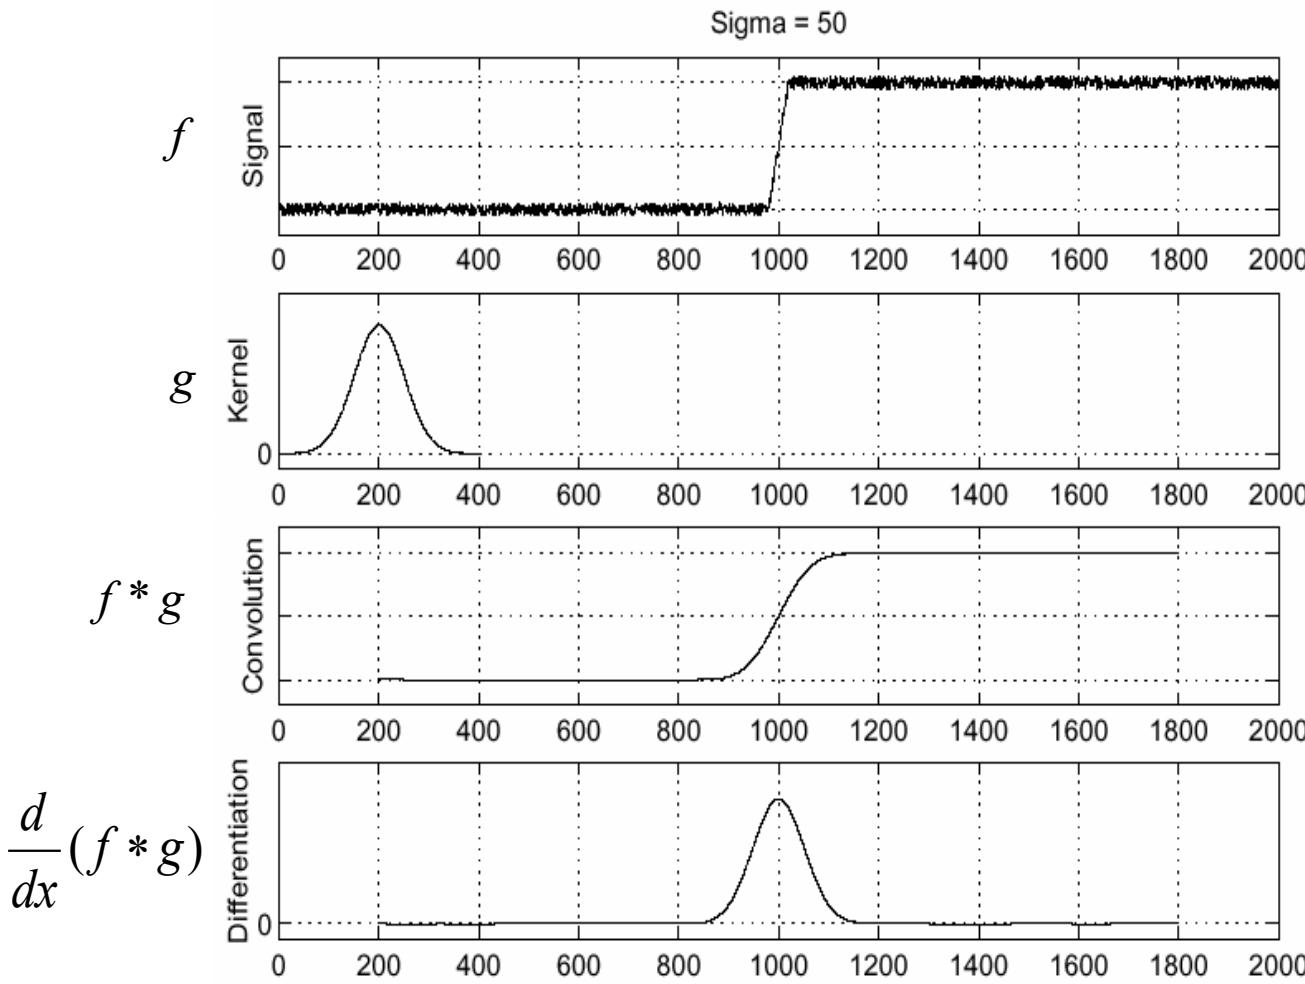
电子科技大学
University of Electronic Science and Technology of China

Source: S. Seitz

Effects of noise

- Difference filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbors
 - Generally, the larger the noise the stronger the response
- What can we do about it?

Solution: smooth first



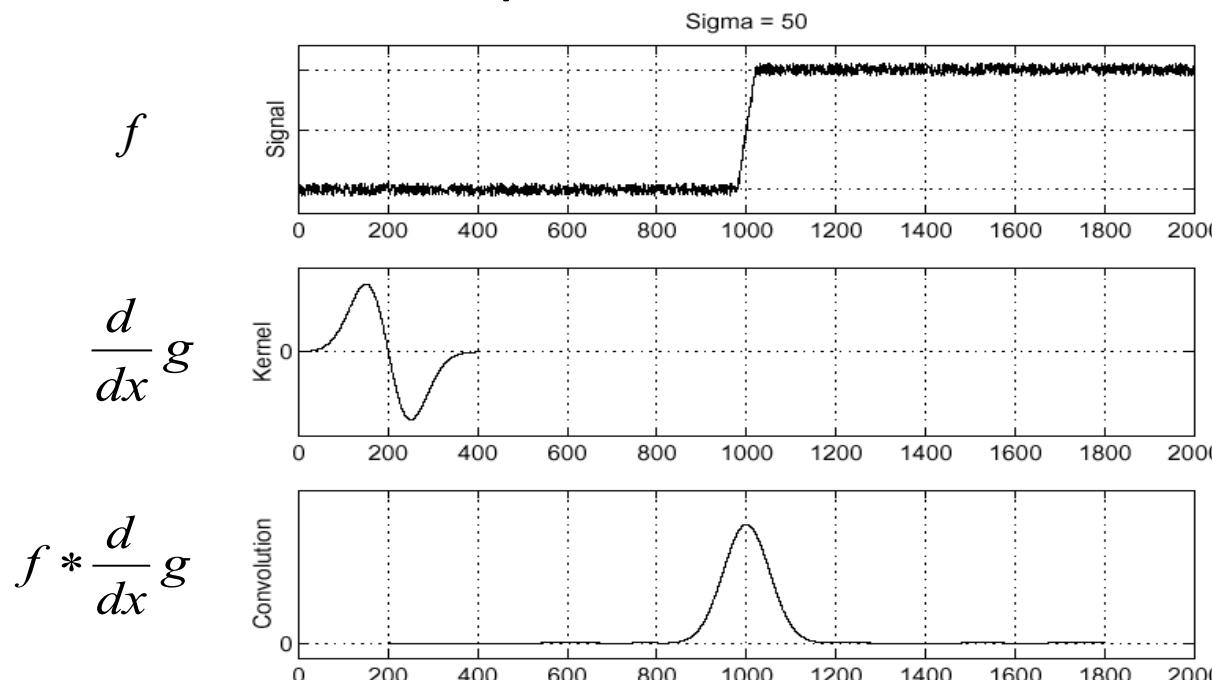
- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

Derivative theorem of convolution

- Convolution is differentiable:

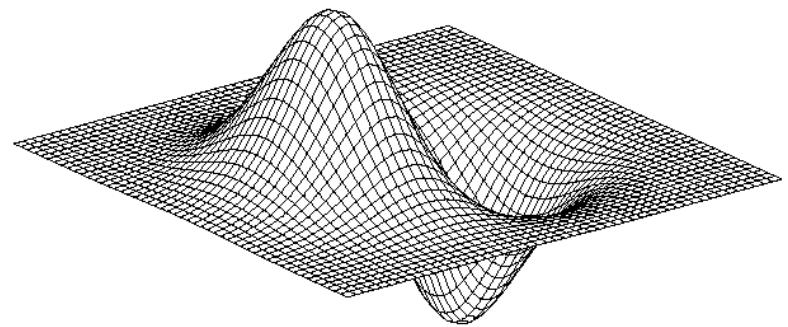
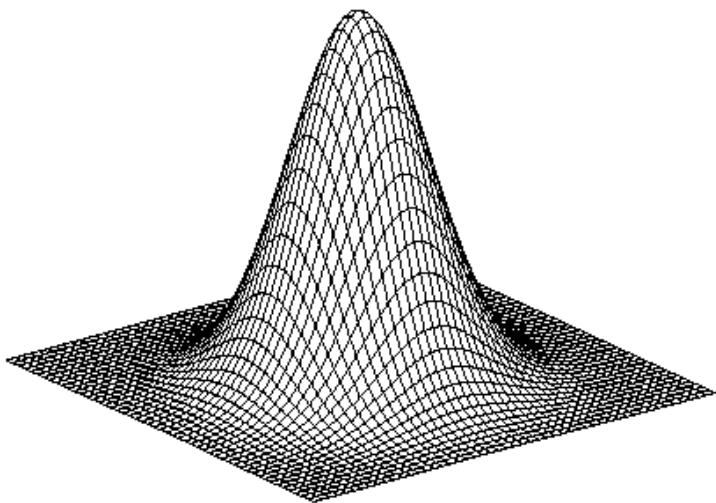
$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

- This saves us one operation:

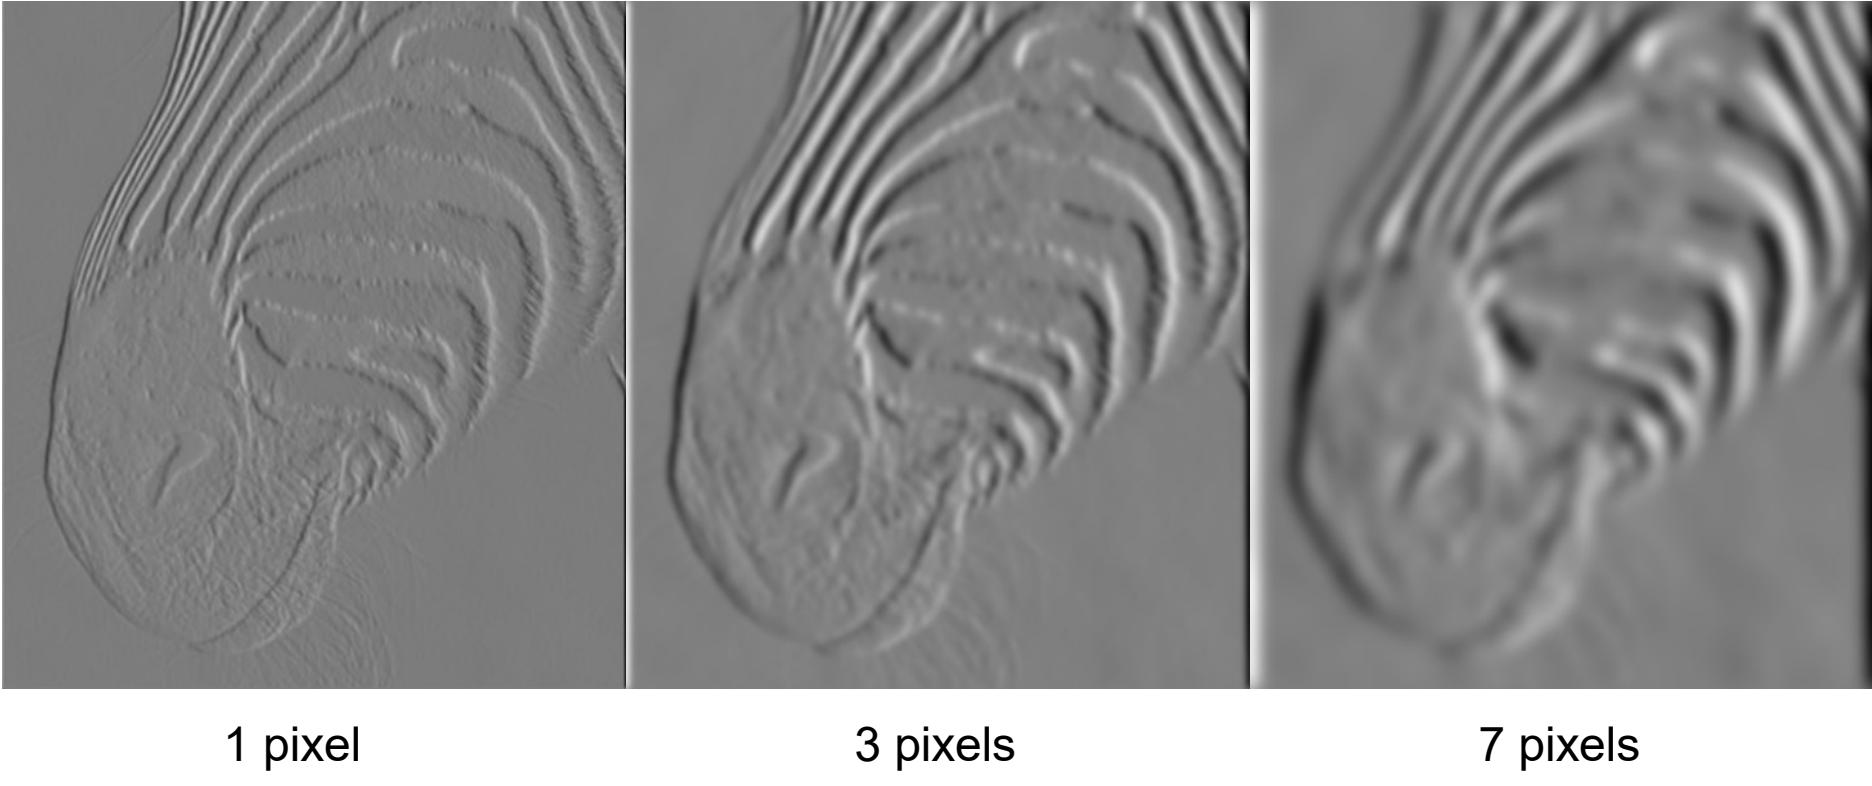


Derivative of 2D Gaussian filter

$$* [1 \ -1] =$$



Tradeoff between smoothing and localization



- Smoothed derivative removes noise, but blurs edge. Also finds edges at different “scales”.

Think-Pair-Share

What is a good edge detector?

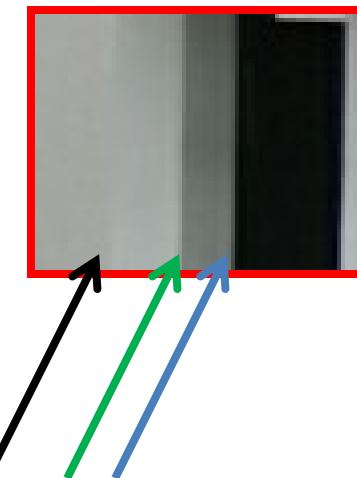
Do we lose information when we look at edges?

Are edges ‘complete’ as a representation of images?

Designing an edge detector

- Criteria for a good edge detector:
 - **Good detection:** the optimal detector should **find all real edges, ignoring noise or other artifacts**
 - **Good localization**
 - the edges detected must be as close as possible to the true edges
 - the detector must return one point only for each true edge point
- Cues of edge detection
 - Differences in color, intensity, or texture across the boundary
 - Continuity and closure
 - High-level knowledge

Closeup of edges



Elder – Are Edges Incomplete? 1999

Edge ‘code’:

- position,
- gradient magnitude,
- gradient direction,
- blur.

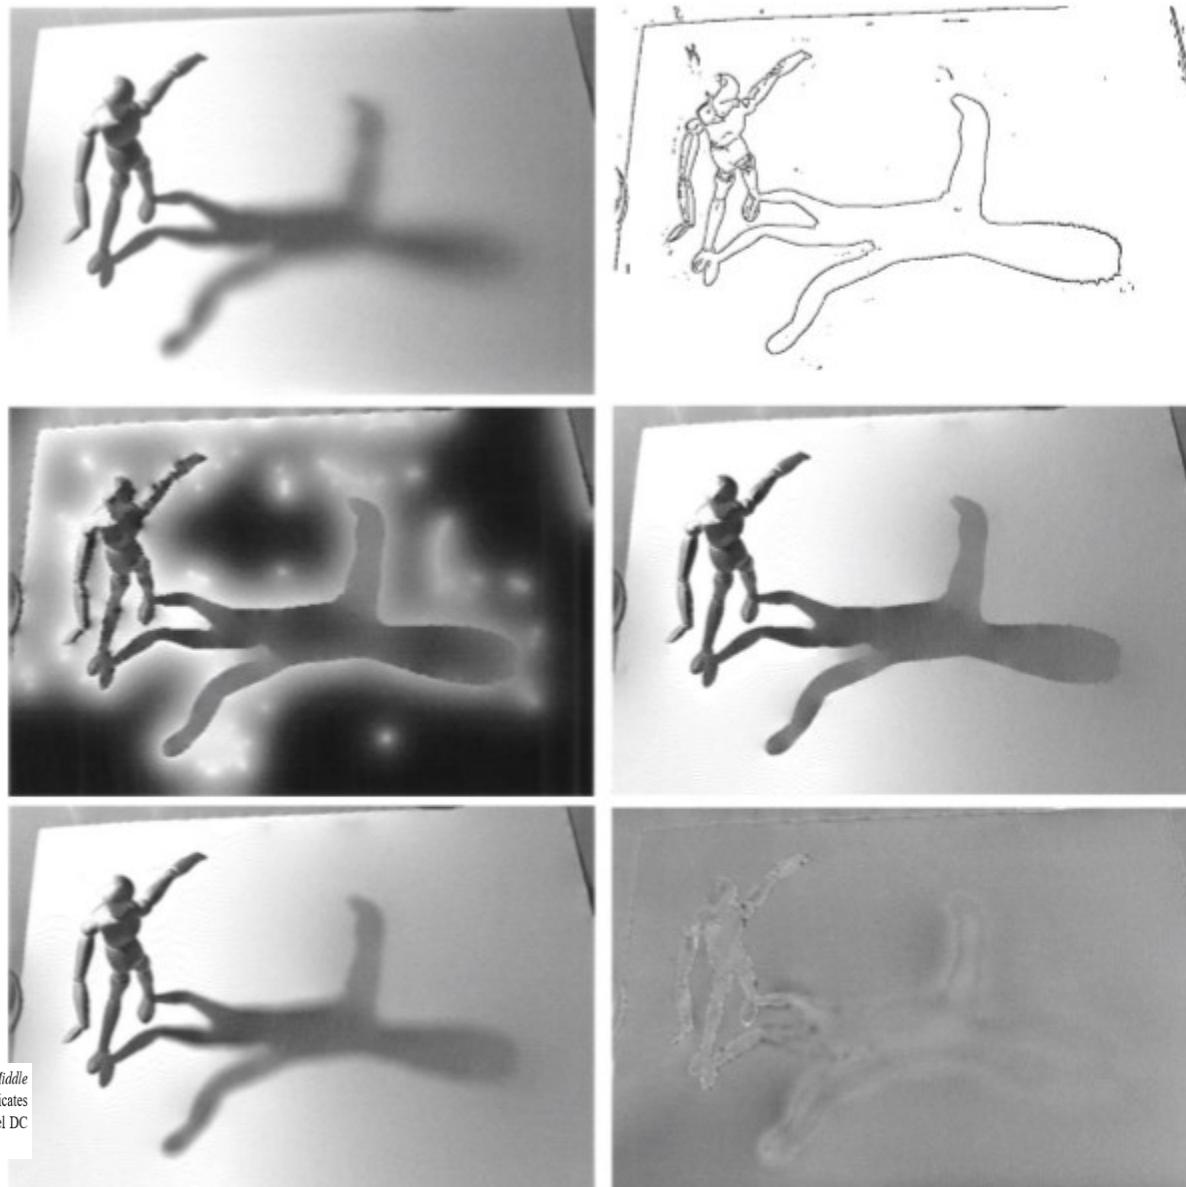


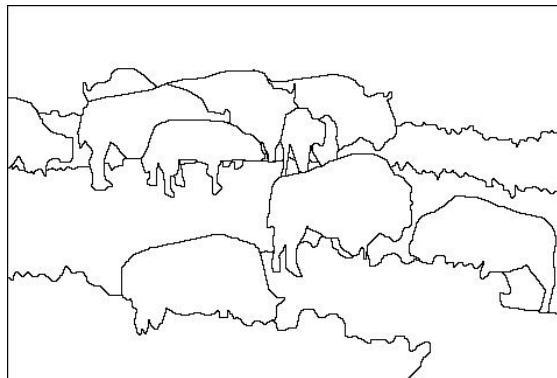
Figure 8. Top left: Original image. Top right: Detected edge locations. Middle left: Intermediate solution to the heat equation. Middle right: Reconstructed luminance function. Bottom left: Reblurred result. Bottom right: Error map (reblurred result—original). Bright indicates overestimation of intensity, dark indicates underestimation. Edge density is 1.7%. RMS error is 10.1 grey levels, with a 3.9 grey level DC component, and an estimated 1.6 grey levels due to noise removal.

Where do humans see boundaries?

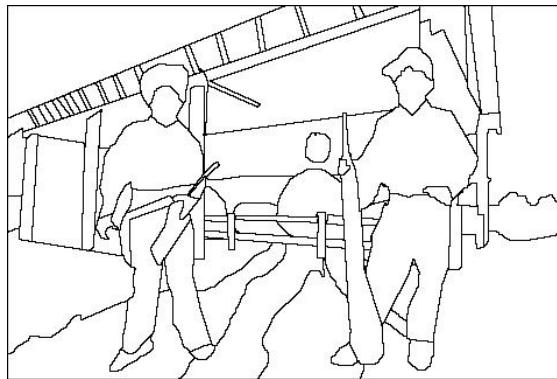
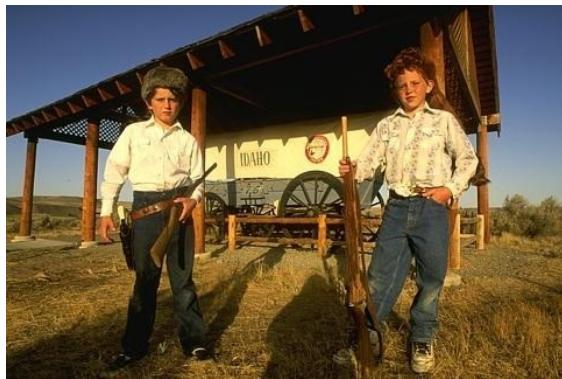
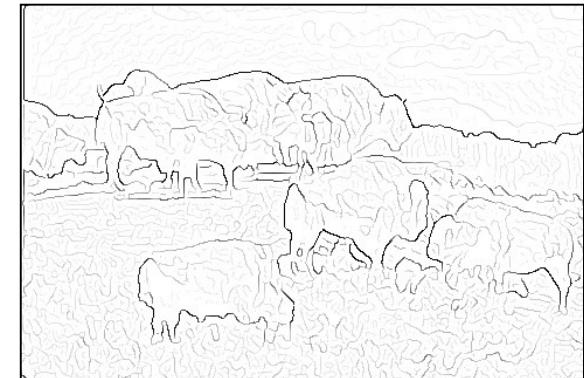
image



human segmentation



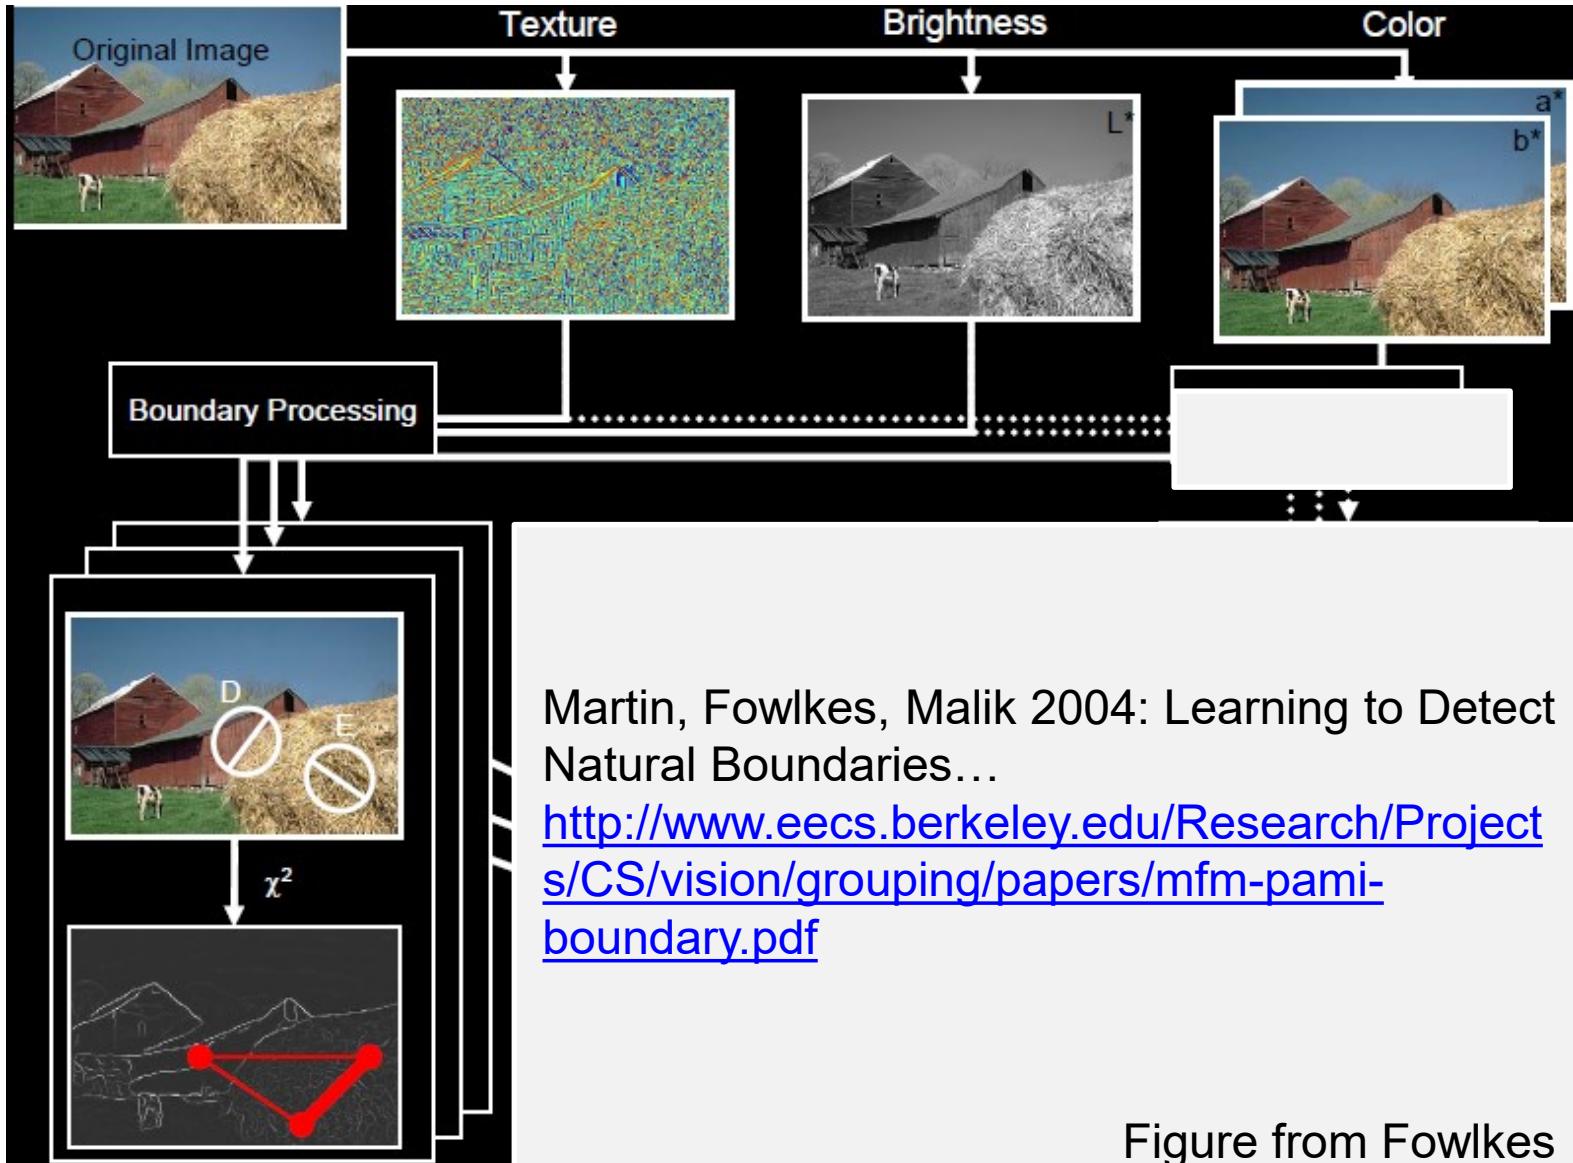
gradient magnitude



- Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

pB boundary detector



未来媒体研究中心
CENTER FOR FUTURE MEDIA



电子科技大学
University of Electronic Science and Technology of China

Brightness

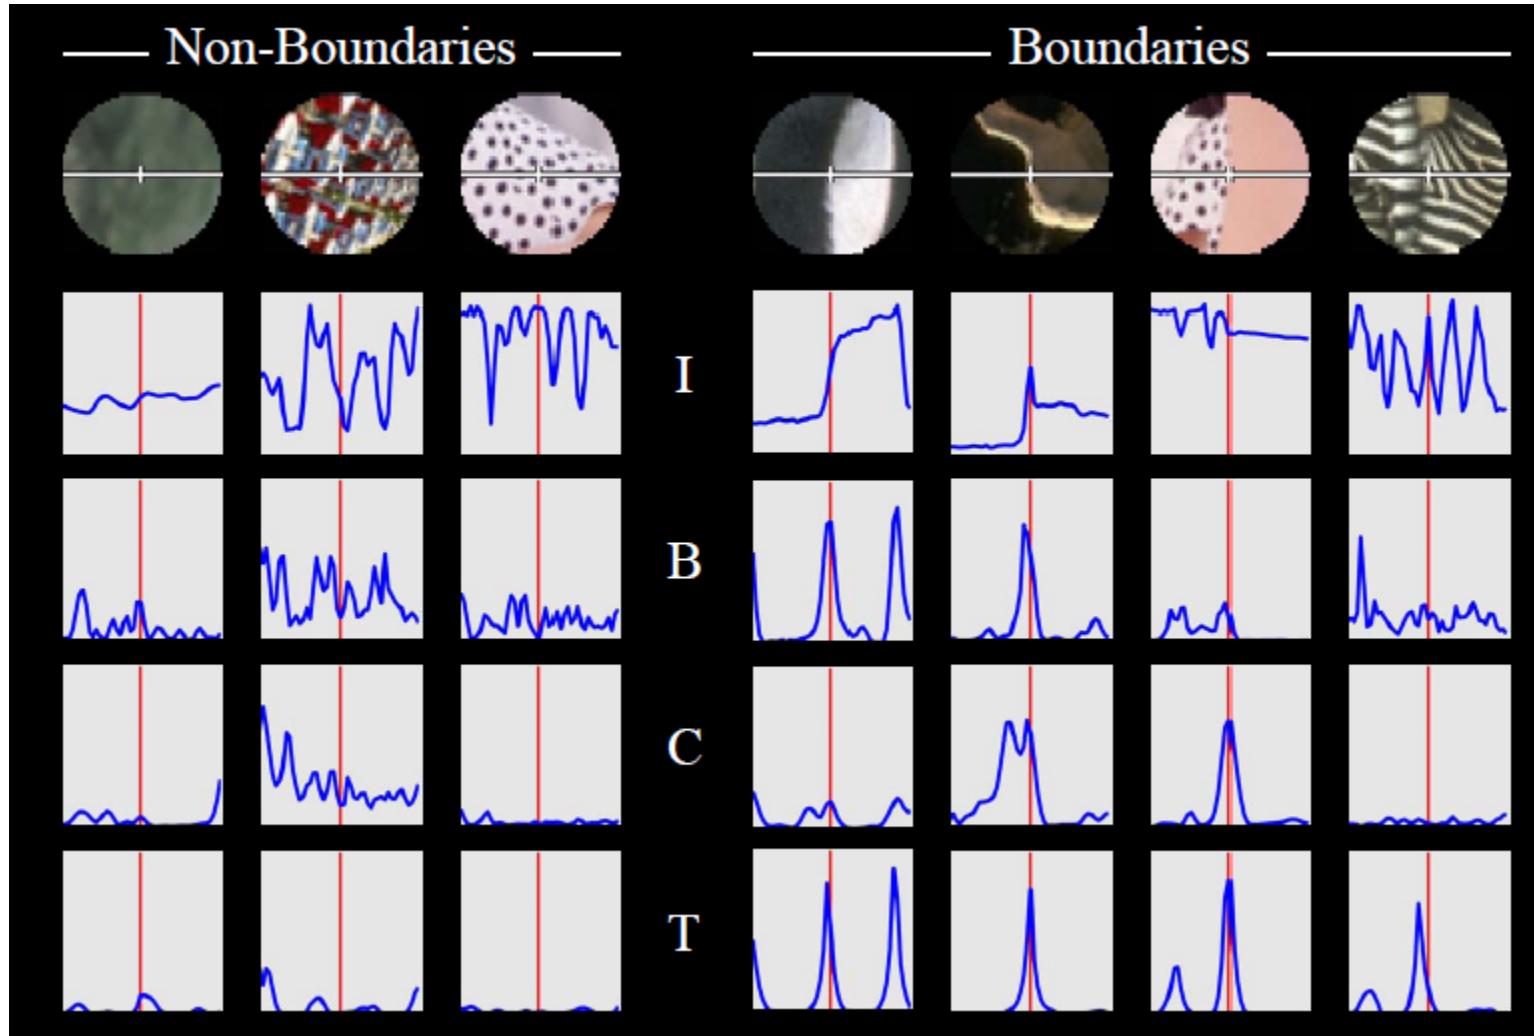


未来媒体研究中心
CENTER FOR FUTURE MEDIA



电子科技大学
University of Electronic Science and Technology of China

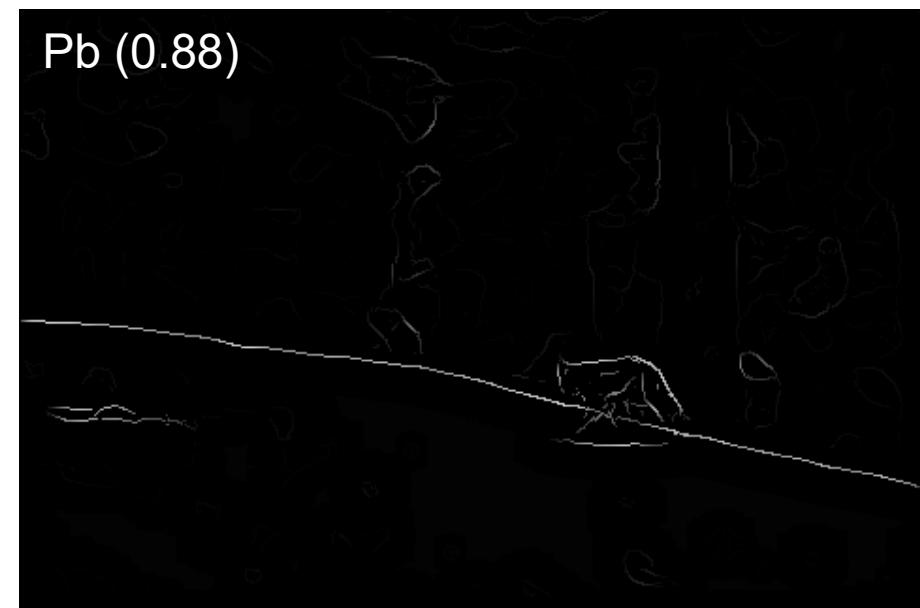
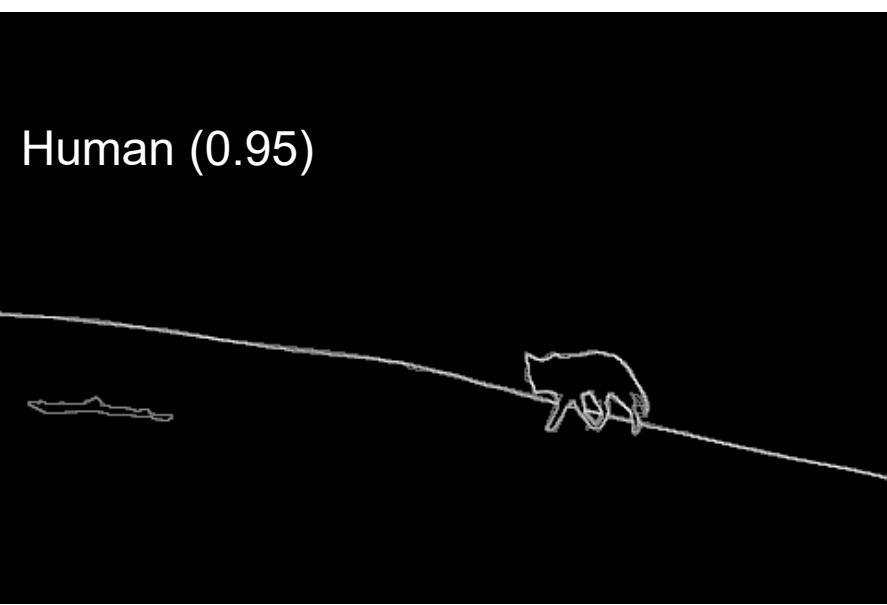
pB Boundary Detector



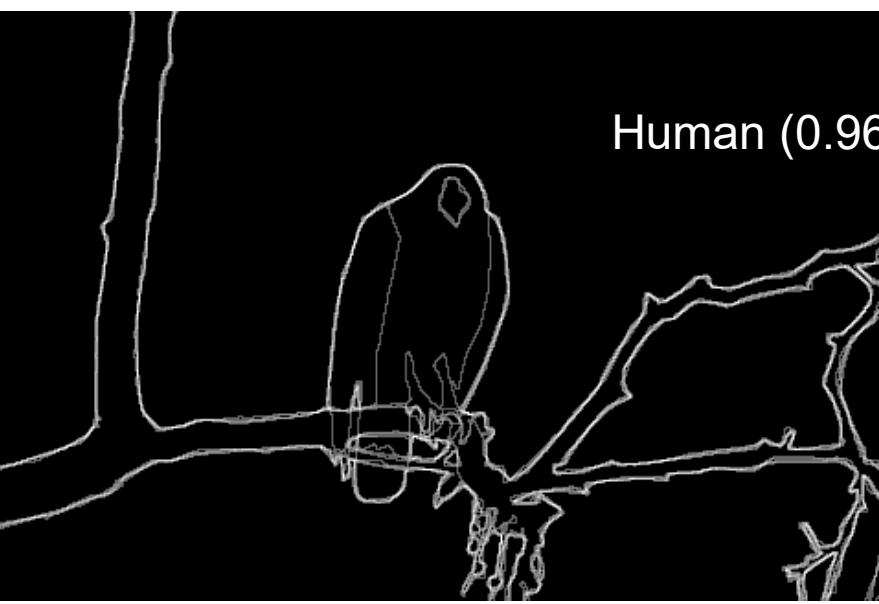
Results



Human (0.95)



Results



Human (0.96)



Pb (0.88)



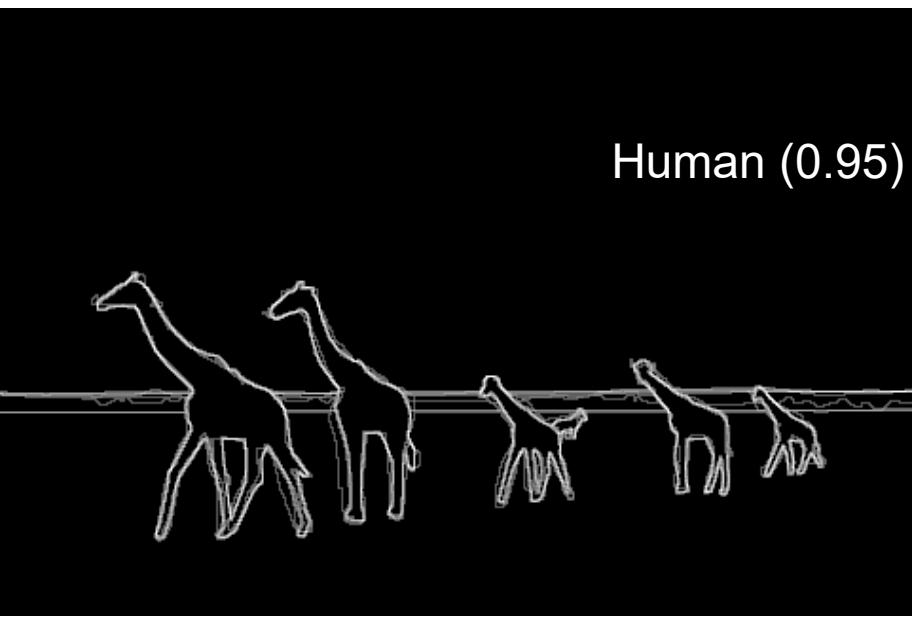
未来媒体研究中心
CENTER FOR FUTURE MEDIA

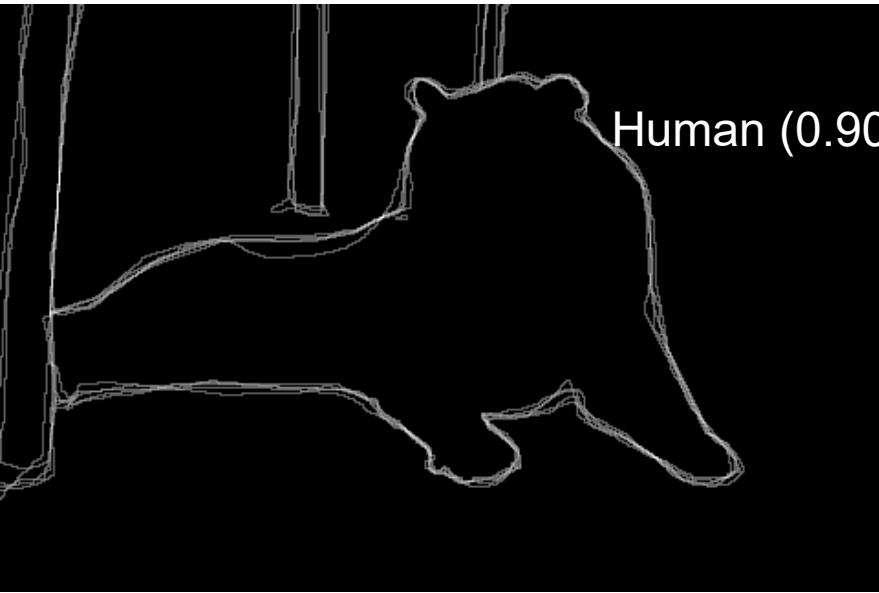


电子科技大学
University of Electronic Science and Technology of China



Human (0.95)





For more:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/bench/html/108082-color.html>

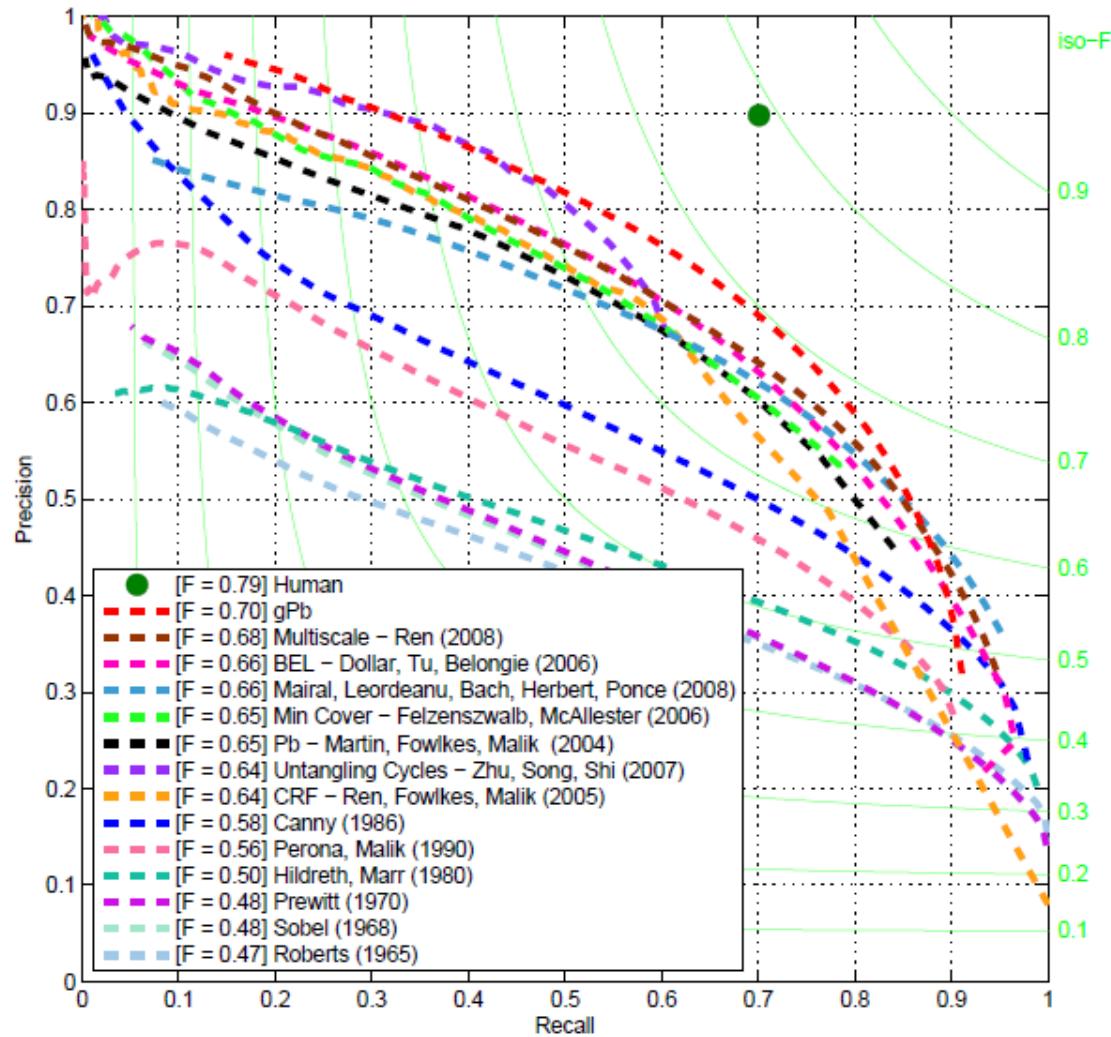


未来媒体研究中心
CENTER FOR FUTURE MEDIA



电子科技大学
University of Electronic Science and Technology of China

45 years of boundary detection

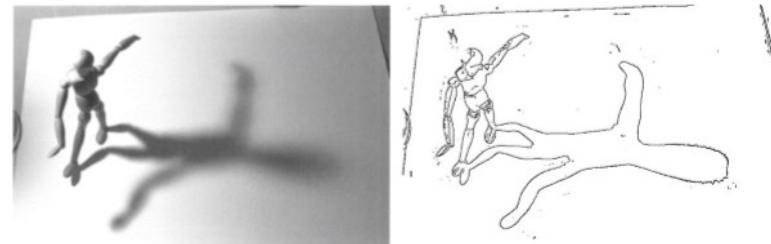
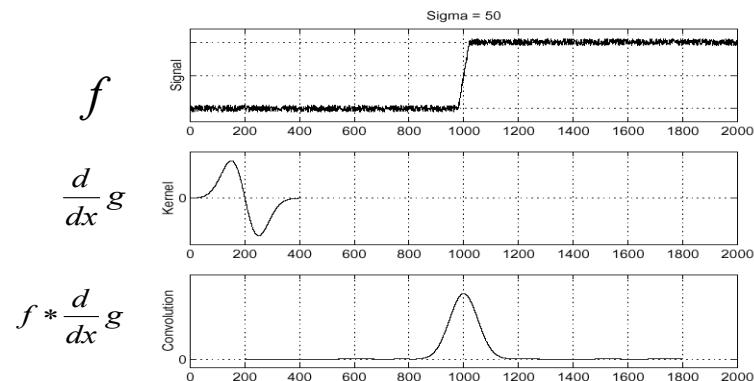


State of edge detection

- Local edge detection works well
 - ‘False positives’ from illumination and texture edges (depends on our application).
- Some methods to take into account longer contours
- Modern methods that actually “learn” from data.
- Poor use of object and high-level information.

Summary: Edges primer

- Edge detection to identify visual change in image
- Derivative of Gaussian and linear combination of convolutions
- What is an edge?
What is a good edge?



Canny edge detector

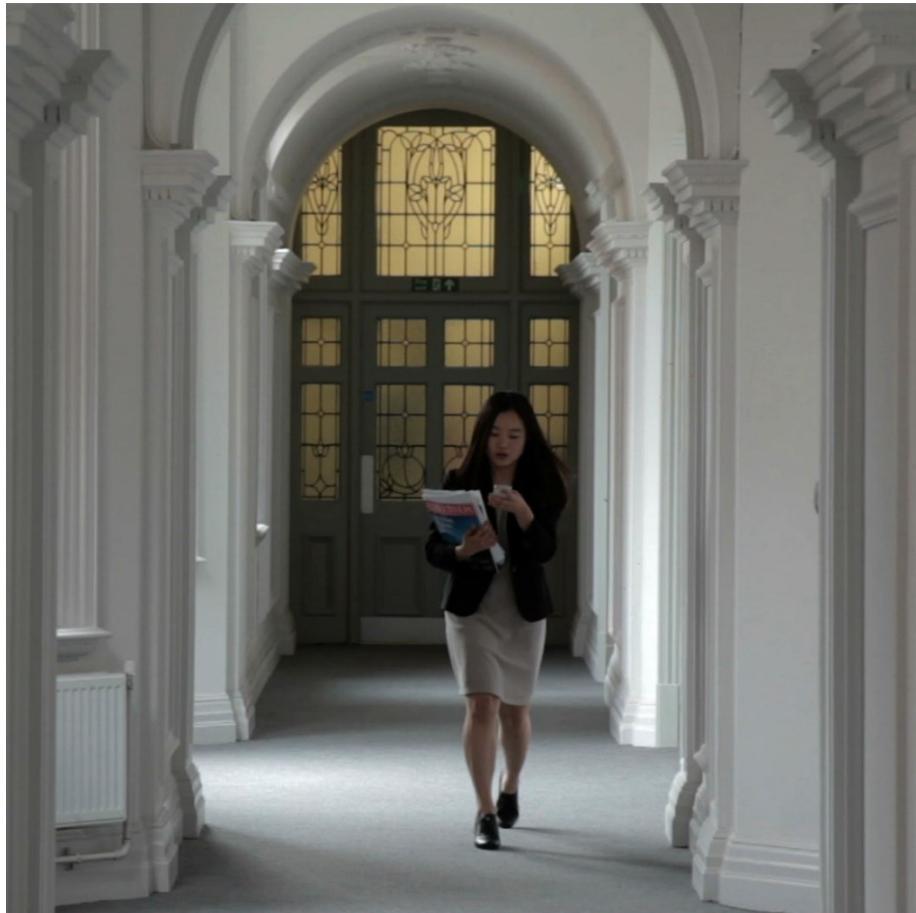
- Probably the most widely used edge detector in computer vision.
- Theoretical model: step-edges corrupted by additive Gaussian noise.
- Canny showed that first derivative of Gaussian closely approximates the operator that optimizes the product of *signal-to-noise ratio* and localization.

J. Canny, [A Computational Approach To Edge Detection](#), IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

22,000 citations!

Demonstrator Image

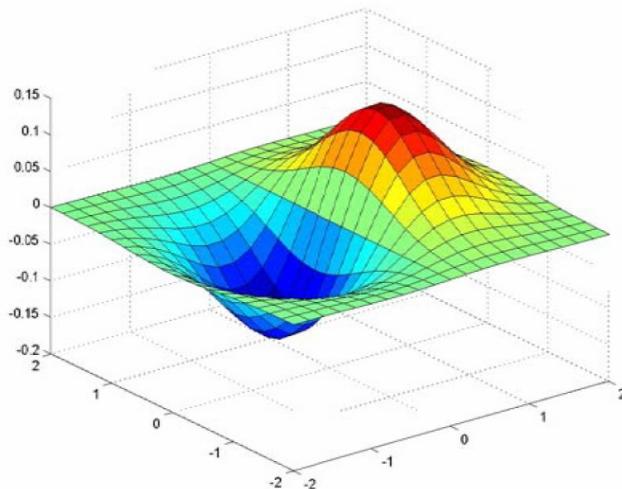
`rgb2gray('img.png')`



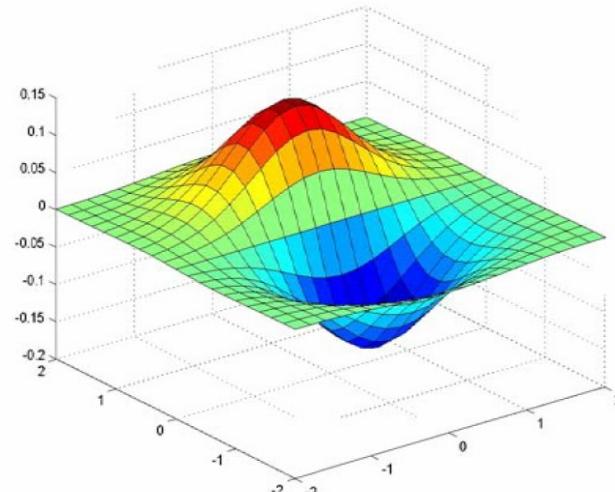
Canny edge detector

1. Filter image with x, y derivatives of Gaussian

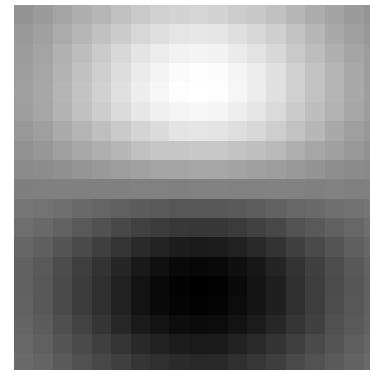
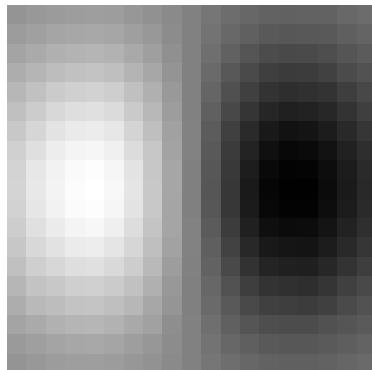
Derivative of Gaussian filter



x-direction



y-direction



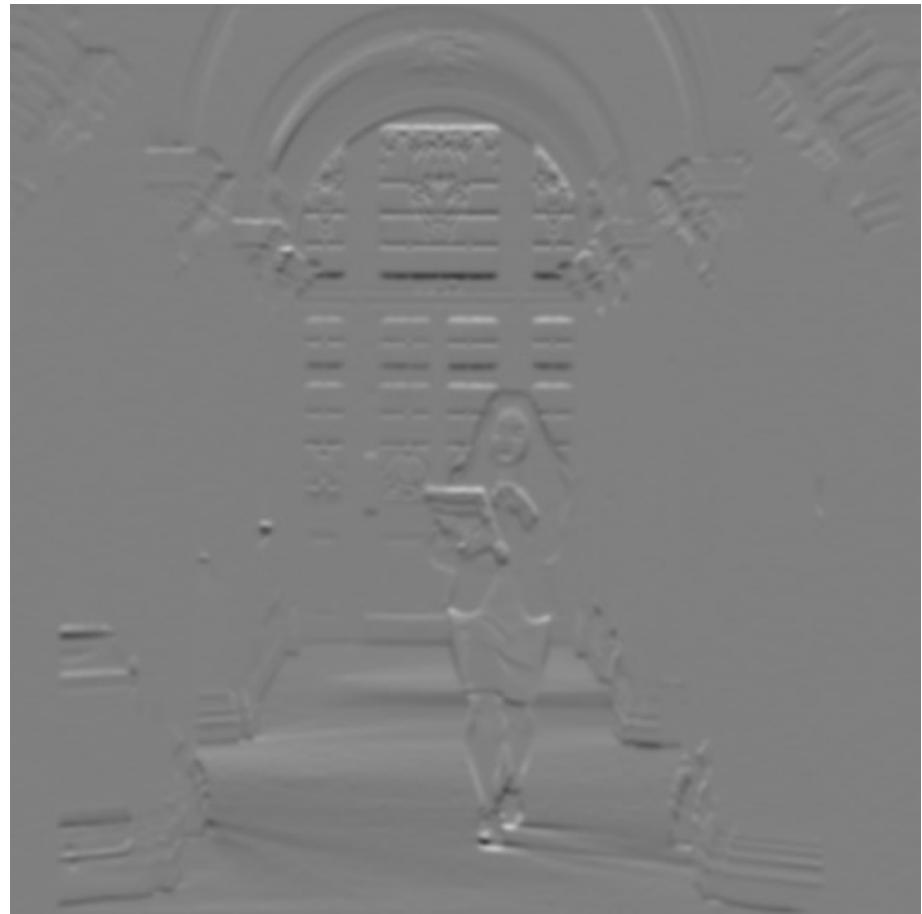
Compute Gradients



X Derivative of Gaussian



Y Derivative of Gaussian



($x2 + 0.5$ for visualization)



未来媒体研究中心
CENTER FOR FUTURE MEDIA



电子科技大学
University of Electronic Science and Technology of China

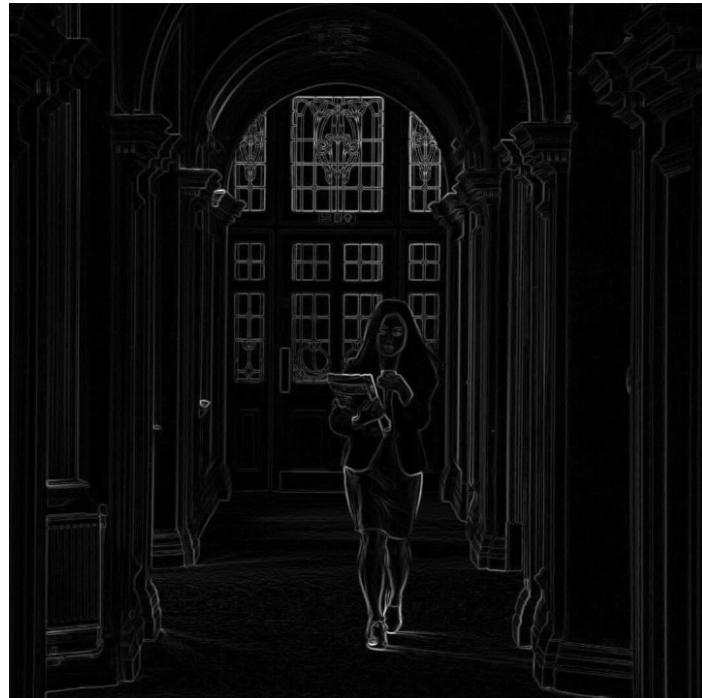
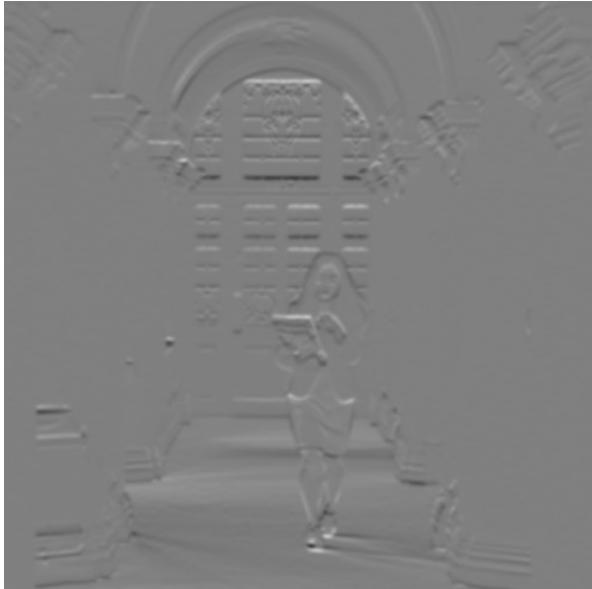
Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient

Compute Gradient Magnitude



$\text{sqrt}(\text{XDerivOfGaussian} .^2 + \text{YDerivOfGaussian} .^2)$ = gradient magnitude



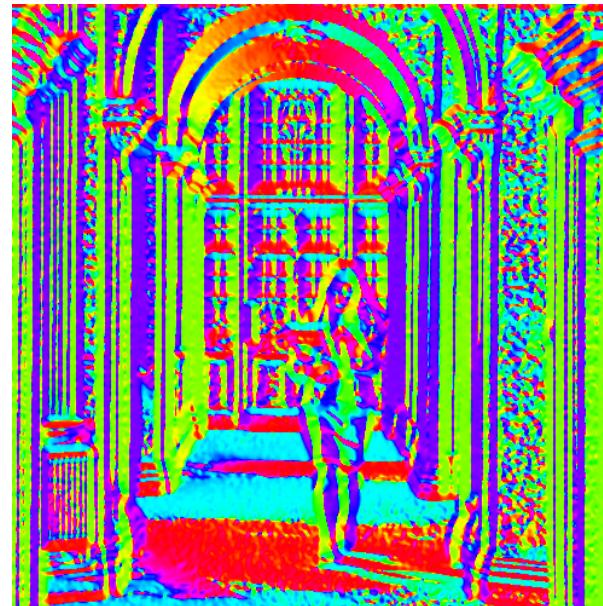
(x4 for visualization)

Compute Gradient Orientation

- Threshold magnitude at minimum level
- Get orientation via $\theta = \text{atan2}(g_y, g_x)$



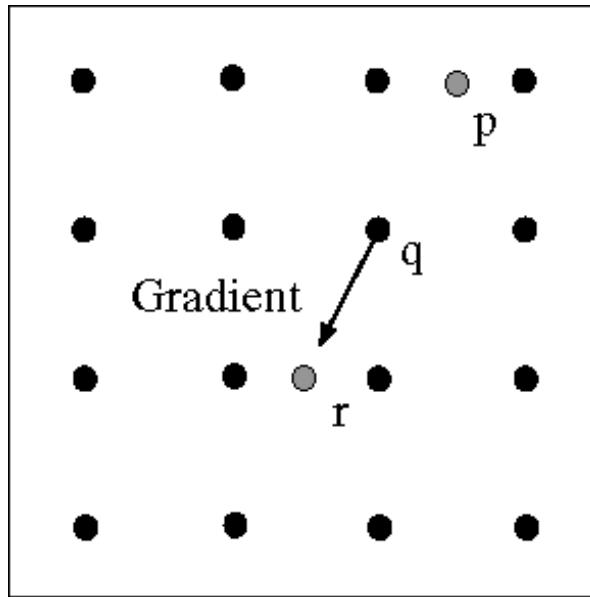
Unthresholded:



Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” to single pixel width

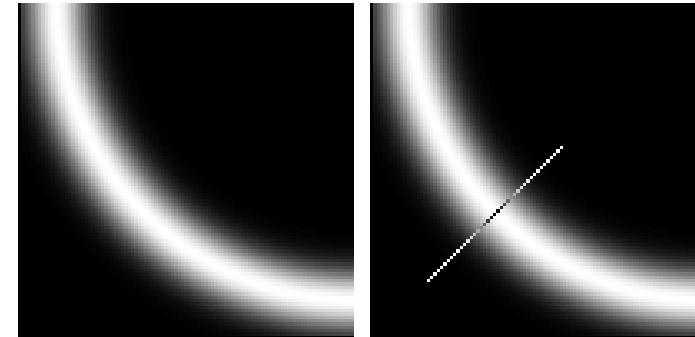
Non-maximum suppression for each orientation



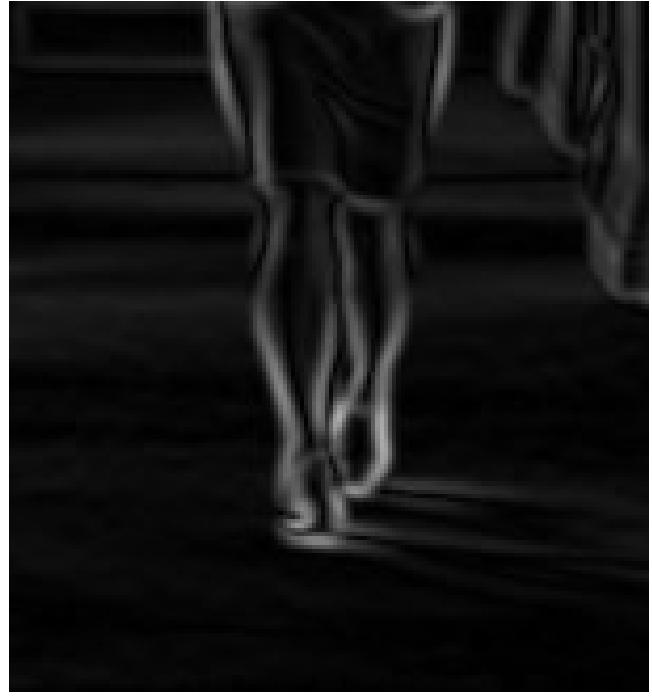
At pixel q:

We have a maximum if the value is larger than those at both p and at r.

Interpolate along gradient direction to get these values.



Before Non-max Suppression



Gradient magnitude (x4 for visualization)

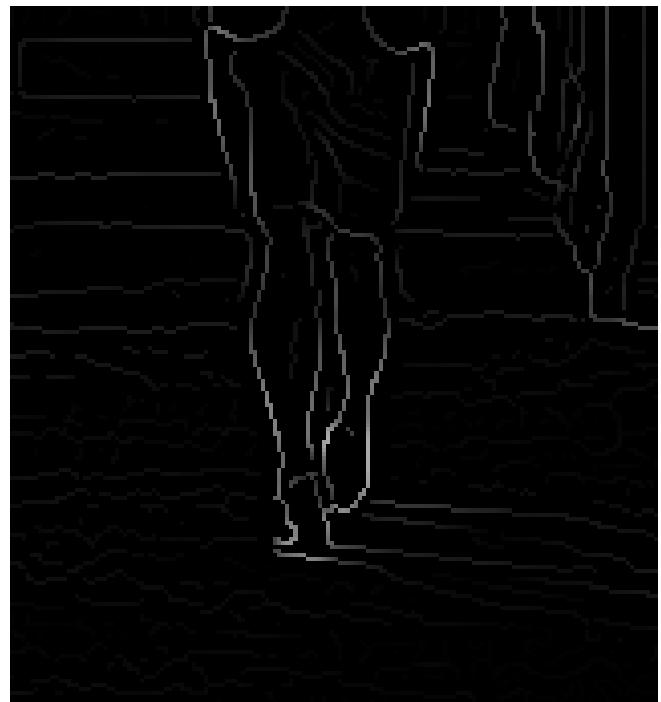


未来媒体研究中心
CENTER FOR FUTURE MEDIA



电子科技大学
University of Electronic Science and Technology of China

After non-max suppression



Gradient magnitude (x4 for visualization)



未来媒体研究中心
CENTER FOR FUTURE MEDIA



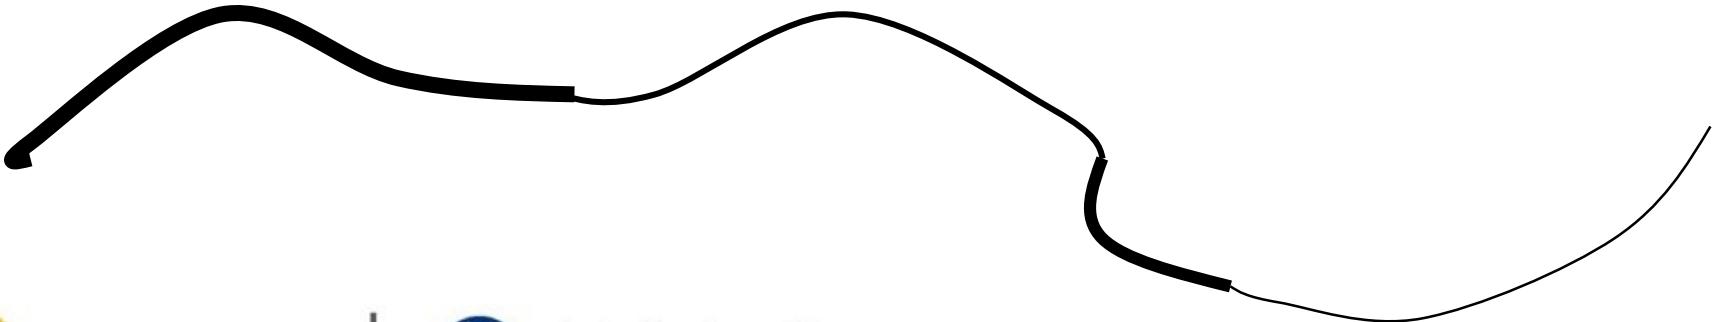
电子科技大学
University of Electronic Science and Technology of China

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” to single pixel width
4. ‘Hysteresis’ Thresholding

'Hysteresis' thresholding

- Two thresholds – high and low
 - Grad. mag. > high threshold? = strong edge
 - Grad. mag. < low threshold? noise
 - In between = weak edge
-
- ‘Follow’ edges starting from strong edge pixels
 - Continue them into weak edges
 - Connected components (Szeliski 3.3.4)



Final Canny Edges

$$\sigma = \sqrt{2}, t_{low} = 0.05, t_{high} = 0.1$$

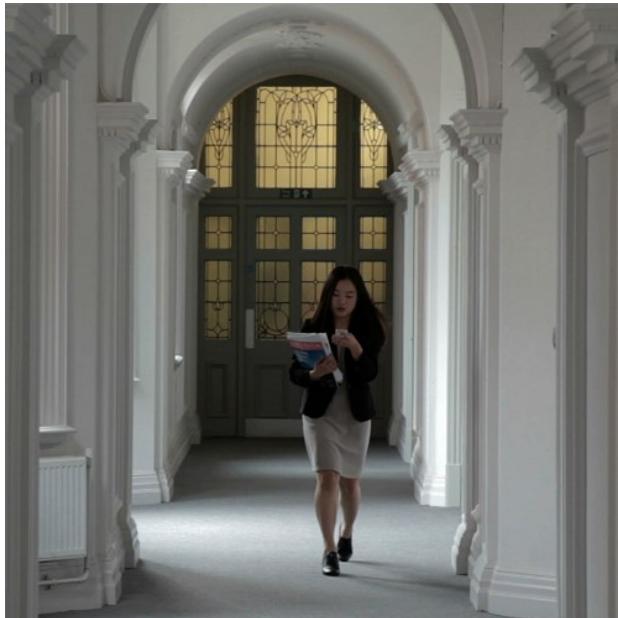


未来媒体研究中心
CENTER FOR FUTURE MEDIA

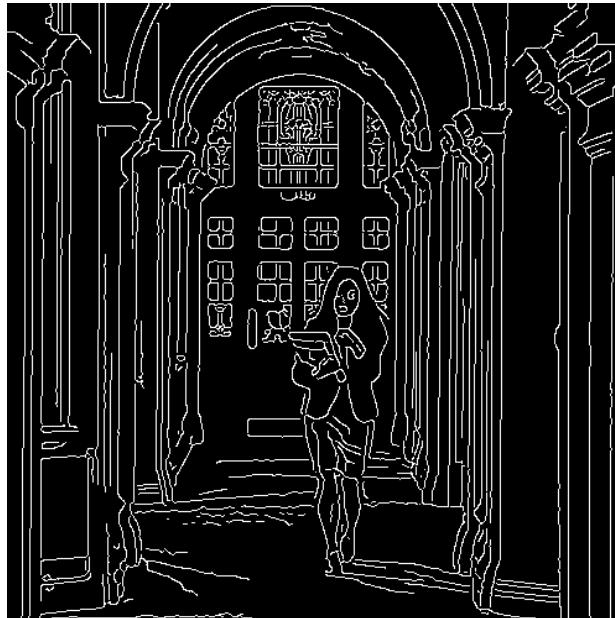


电子科技大学
University of Electronic Science and Technology of China

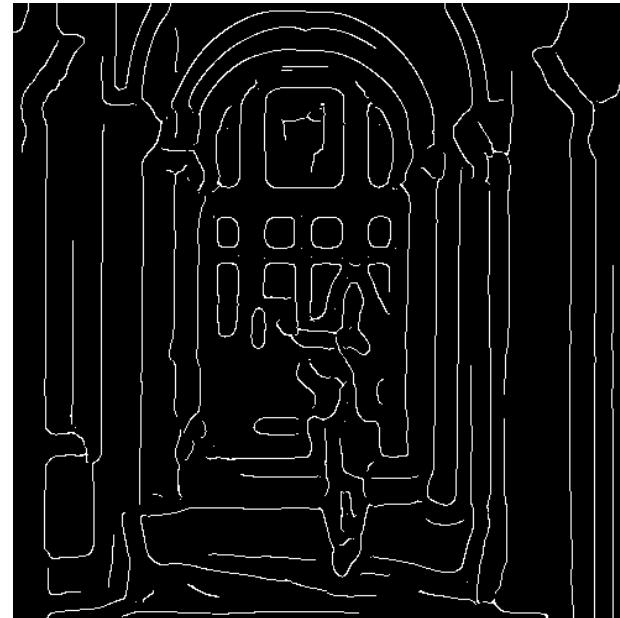
Effect of σ (Gaussian kernel spread/size)



Original



$\sigma = \sqrt{2}$



$\sigma = 4\sqrt{2}$

The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features

Canny edge detector

1. Filter image with x, y derivatives of Gaussian
 2. Find magnitude and orientation of gradient
 3. Non-maximum suppression:
 - Thin multi-pixel wide “ridges” to single pixel width
 4. ‘Hysteresis’ Thresholding:
 - Define two thresholds: low and high
 - Use the high threshold to start edge curves and the low threshold to continue them
 - ‘Follow’ edges starting from strong edge pixels
 - Connected components (Szeliski 3.3.4)
-
- MATLAB: `edge(image, 'canny')`